

An Introduction to SparkR

The Apache Spark Framework

Roland Boubela

Center for Medical Physics and Biomedical Engineering
Medical University of Vienna

26. Jan. 2016



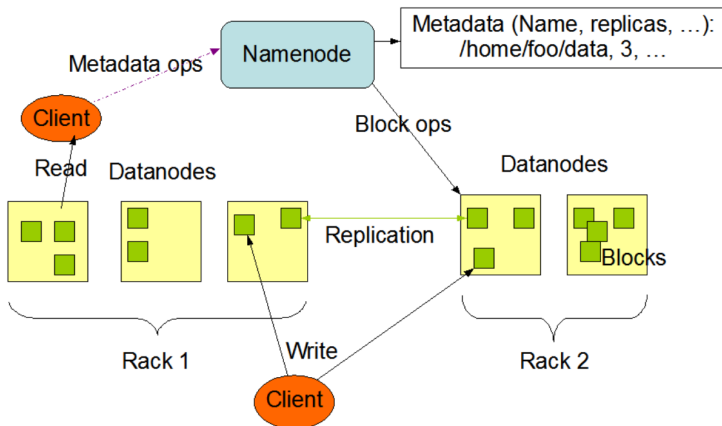
Hadoop Distributed Filesystem (HDFS)



- ▶ Based on Google filesystem
- ▶ Fault tolerant (distributed blocks, multiple replicas)
- ▶ One client can write, multiple clients read
- ▶ Only sequential r/w-operations (no random access)

Hadoop Distributed Filesystem (HDFS)

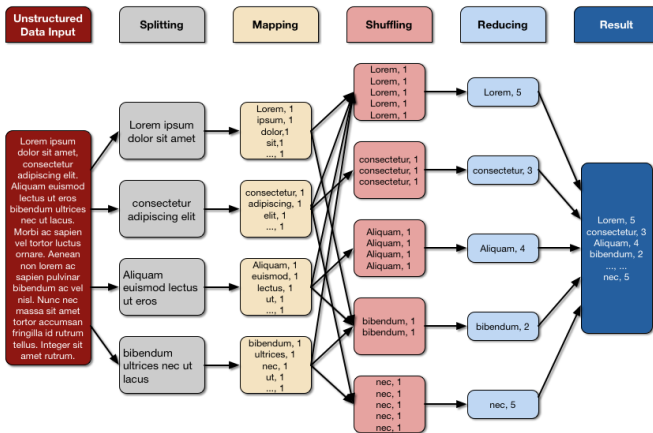
HDFS Architecture



MapReduce Paradigm

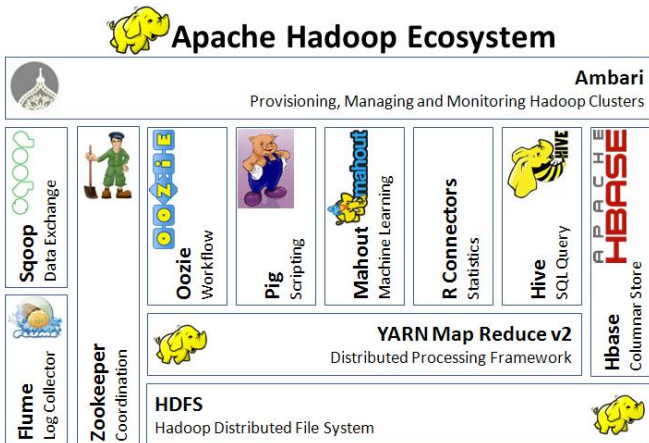
- ▶ MapReduce: Simplified Data Processing on Large Clusters (Dean and Gemawat 2004, OSDI)
- ▶ Splits computations in *map* and *reduce* phase
- ▶ Handles
 - ▶ Details of input data partitioning
 - ▶ Scheduling program's execution across a set of machines
 - ▶ Machine failures
 - ▶ Required inter-machine communication

MapReduce Data and Process Flow of Word Count



(c) 2014 Broomfield Technology Consultants. All Rights Reserved.

Hadoop Ecosystem



Matei Zaharia et al. Paper

Resilient Distributed Datasets: A Fault-Tolerant Abstraction for In-Memory Cluster Computing

Matei Zaharia, Mosharaf Chowdhury, Tathagata Das, Ankur Dave, Justin Ma,
Murphy McCauley, Michael J. Franklin, Scott Shenker, Ion Stoica
University of California, Berkeley

Spark Concept




Figure 2: Spark runtime. The user's driver program launches multiple workers, which read data blocks from a distributed file system and can persist computed RDD partitions in memory.

RDD Properties

- ▶ Resilient, distributed collections
- ▶ Immutable
- ▶ Transformations
 - ▶ map, filter, reduceByKey, join, ...
- ▶ Actions
 - ▶ reduce, collect, count, foreach, ...



RDD Example

```
scala> val textFile = sc.textFile("README.md")
textFile: spark.RDD[String] = spark.MappedRDD@2ee9b6e3


scala> textFile.count() // Number of items in this RDD
res0: Long = 126

scala> textFile.first() // First item in this RDD
res1: String = # Apache Spark

scala> val linesWithSpark = textFile.filter(line => line.contains("Spark"))
linesWithSpark: spark.RDD[String] = spark.FilteredRDD@7dd4af09

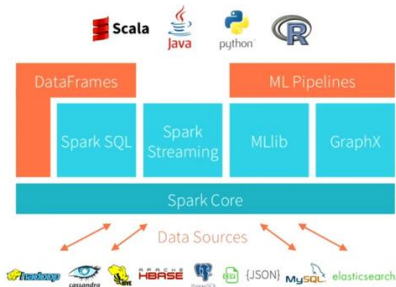
scala> textFile.filter(line => line.contains("Spark")).count()
res3: Long = 15
```

Spark Libraries



Current Situation

- ▶ Current version 1.6.0 (January 2016)
- ▶ SparkR included since 1.3.0 (March 2015)



What is SparkR

- ▶ Light-weight frontend to Apache Spark from R
- ▶ SparkR's development is very *dynamic*
- ▶ Main feature: SparkR DataFrames
 - ▶ distributed collection of data organized into named columns
 - ▶ conceptually equivalent to
 - ▶ a table in a relational database or
 - ▶ a data frame in R
 - ... but on large datasets
- ▶ GLM using MLib

R Example

```
df <- createDataFrame(sqlContext, faithful)

# Get basic information about the DataFrame
df
## DataFrame[eruptions:double, waiting:double]

# Displays the content of the DataFrame to stdout
head(df)
##   eruptions waiting
##1      3.600      79
##2      1.800      54
##3      3.333      74
```

R Example

```
# Filter the DataFrame to only retain rows with wait times shorter than 50 mins
head(filter(df, df$waiting < 50))
## eruptions waiting
##1      1.750      47
##2      1.750      47
##3      1.867      48

# We use the 'n' operator to count the number of times each waiting time appears
head(summarize(groupBy(df, df$waiting), count = n(df$waiting)))
## waiting count
##1      81      13
##2      60       6
##3      68       1

# We can also sort the output from the aggregation to get the most common waiting times
waiting_counts <- summarize(groupBy(df, df$waiting), count = n(df$waiting))
head(arrange(waiting_counts, desc(waiting_counts$count)))

## waiting count
##1      78      15
##2      83      14
##3      81      13
```

Reading from Data Sources

```
sc <- sparkR.init(sparkPackages="com.databricks:spark-csv_2.11:1.0.3")
sqlContext <- sparkRSQL.init(sc)

people <- read.df(sqlContext, "./examples/src/main/resources/people.json", "json")
head(people)
##   age   name
##1  NA Michael
##2  30    Andy
##3  19   Justin

# SparkR automatically infers the schema from the JSON file
printSchema(people)
# root
# |-- age: integer (nullable = true)
# |-- name: string (nullable = true)

write.df(people, path="people.parquet", source="parquet", mode="overwrite")
```


Cluster setup

- ▶ Apache Spark in standalone mode
- ▶ 5 Servers
 - ▶ 1 master node
 - ▶ 4 worker nodes
- ▶ Node specifications
 - ▶ 2 x 4 core Intel Xeon Processor
 - ▶ 48 GB RAM
 - ▶ 1 GBit network to storage
 - ▶ 20 GBit IP over InfiniBand (interconnect)

Connecting to the Apache Spark Master

```
./bin/sparkR --master spark://10.110.100.120:7077 \  
--packages com.databricks:spark-csv_2.10:1.3.0
```

```
R version 3.1.1 (2014-07-10) -- "Sock it to Me"  
Copyright (C) 2014 The R Foundation for Statistical Computing  
Platform: x86_64-unknown-linux-gnu (64-bit)
```

```
...
```